

Personalization

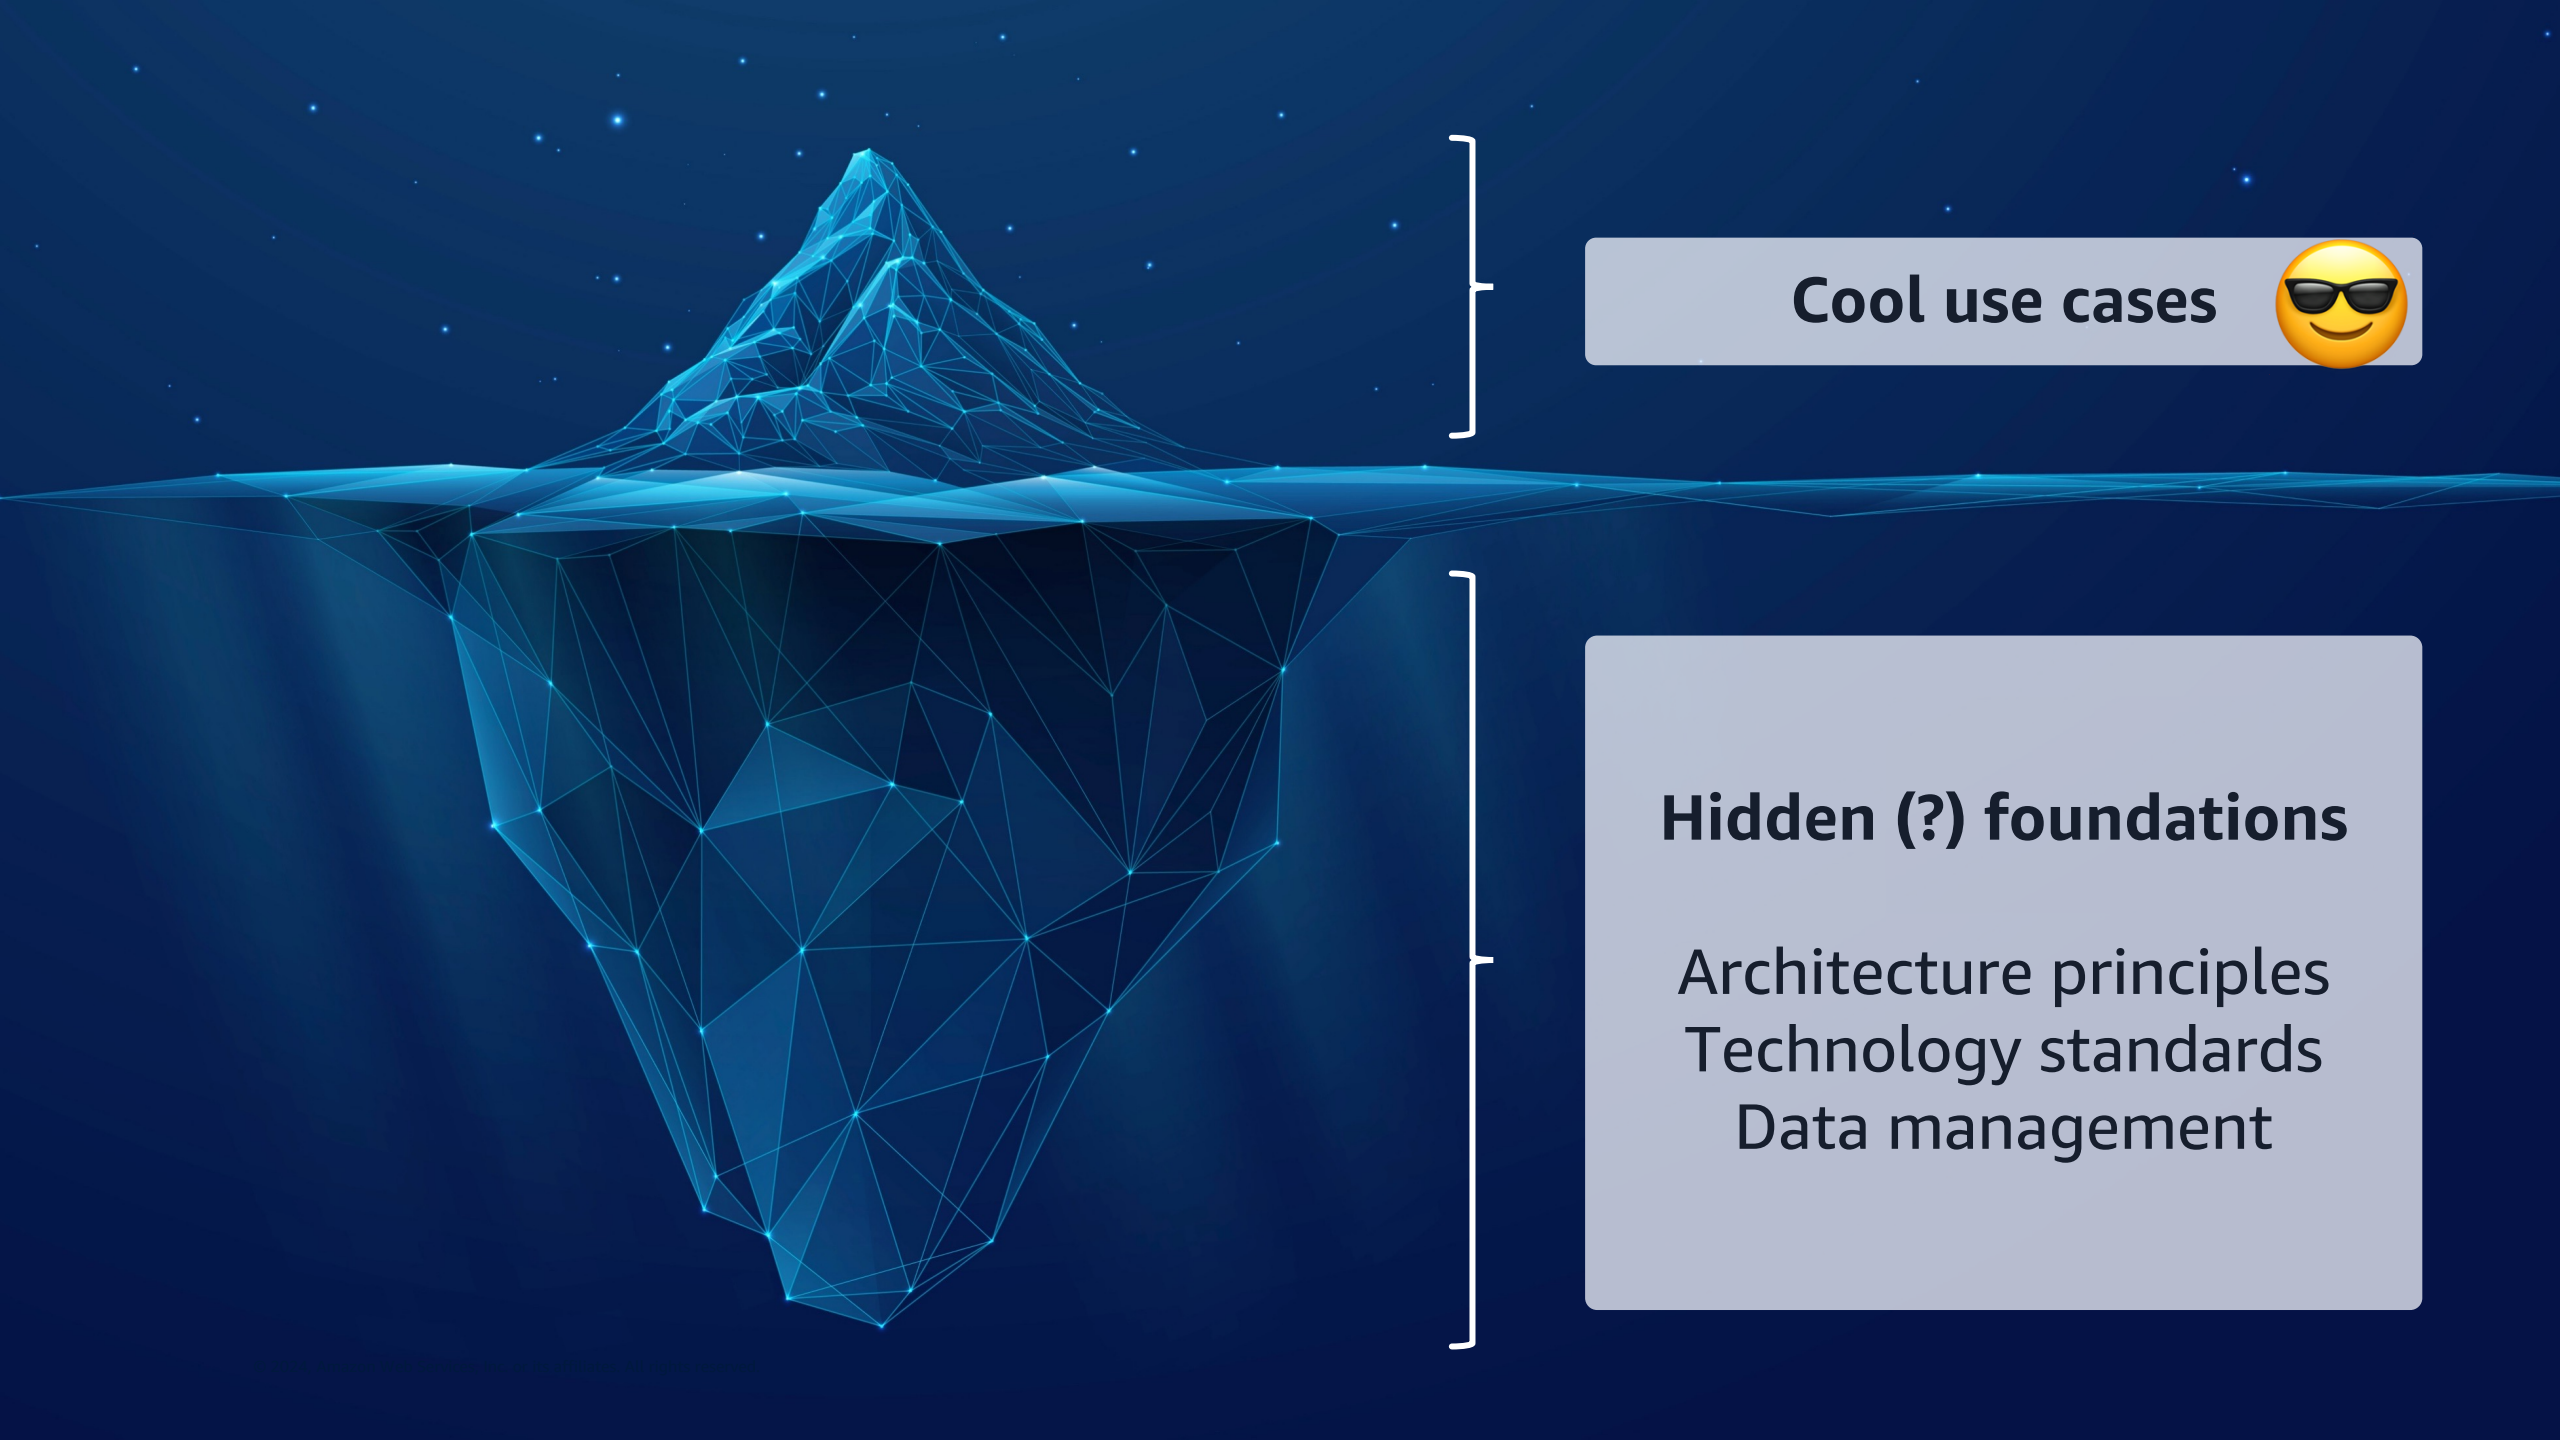
Real-time dashboards

Fraud detection

Clickstream analytics

Real-time logistic
optimization

Generative AI



Cool use cases



Hidden (?) foundations

Architecture principles
Technology standards
Data management

We are here



2024


Data grows **>10x**
every **5 years**

IDC, Data Age 2025



Norman Weisenburger

Techie by heart, connecting people and technology at AWS Retail/CPG customers in Germany as a Senior Solutions Architect at Amazon Web Services

 /in/normalerweise

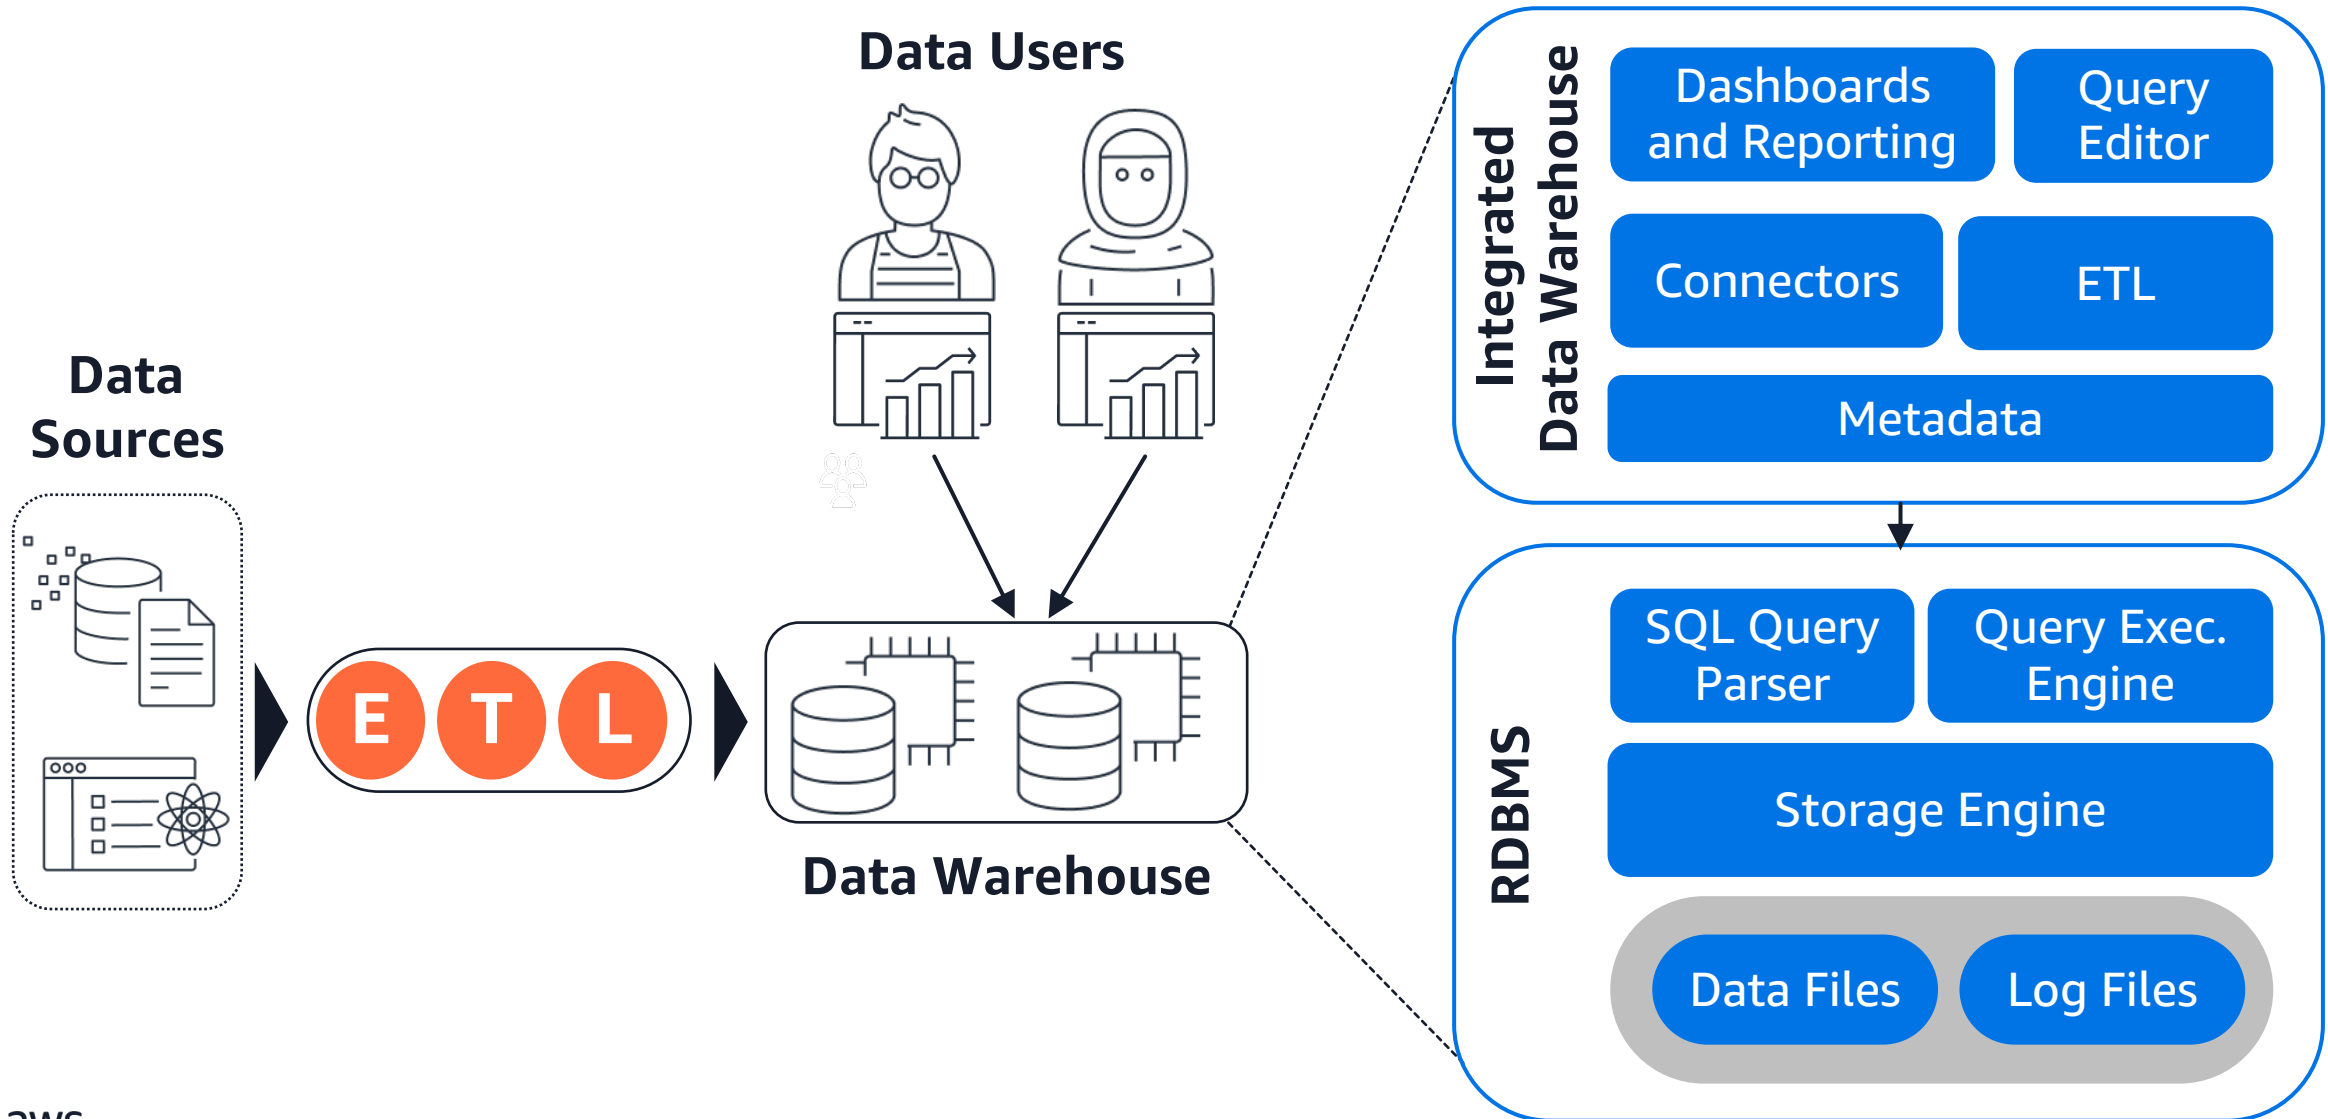


Beyond Warehouses: Transactional Data Lakes Power Modern Analytics

Norman Weisenburger
(he/him)


From Data Warehouses to Data Lakes

The Data Warehouse ~2006

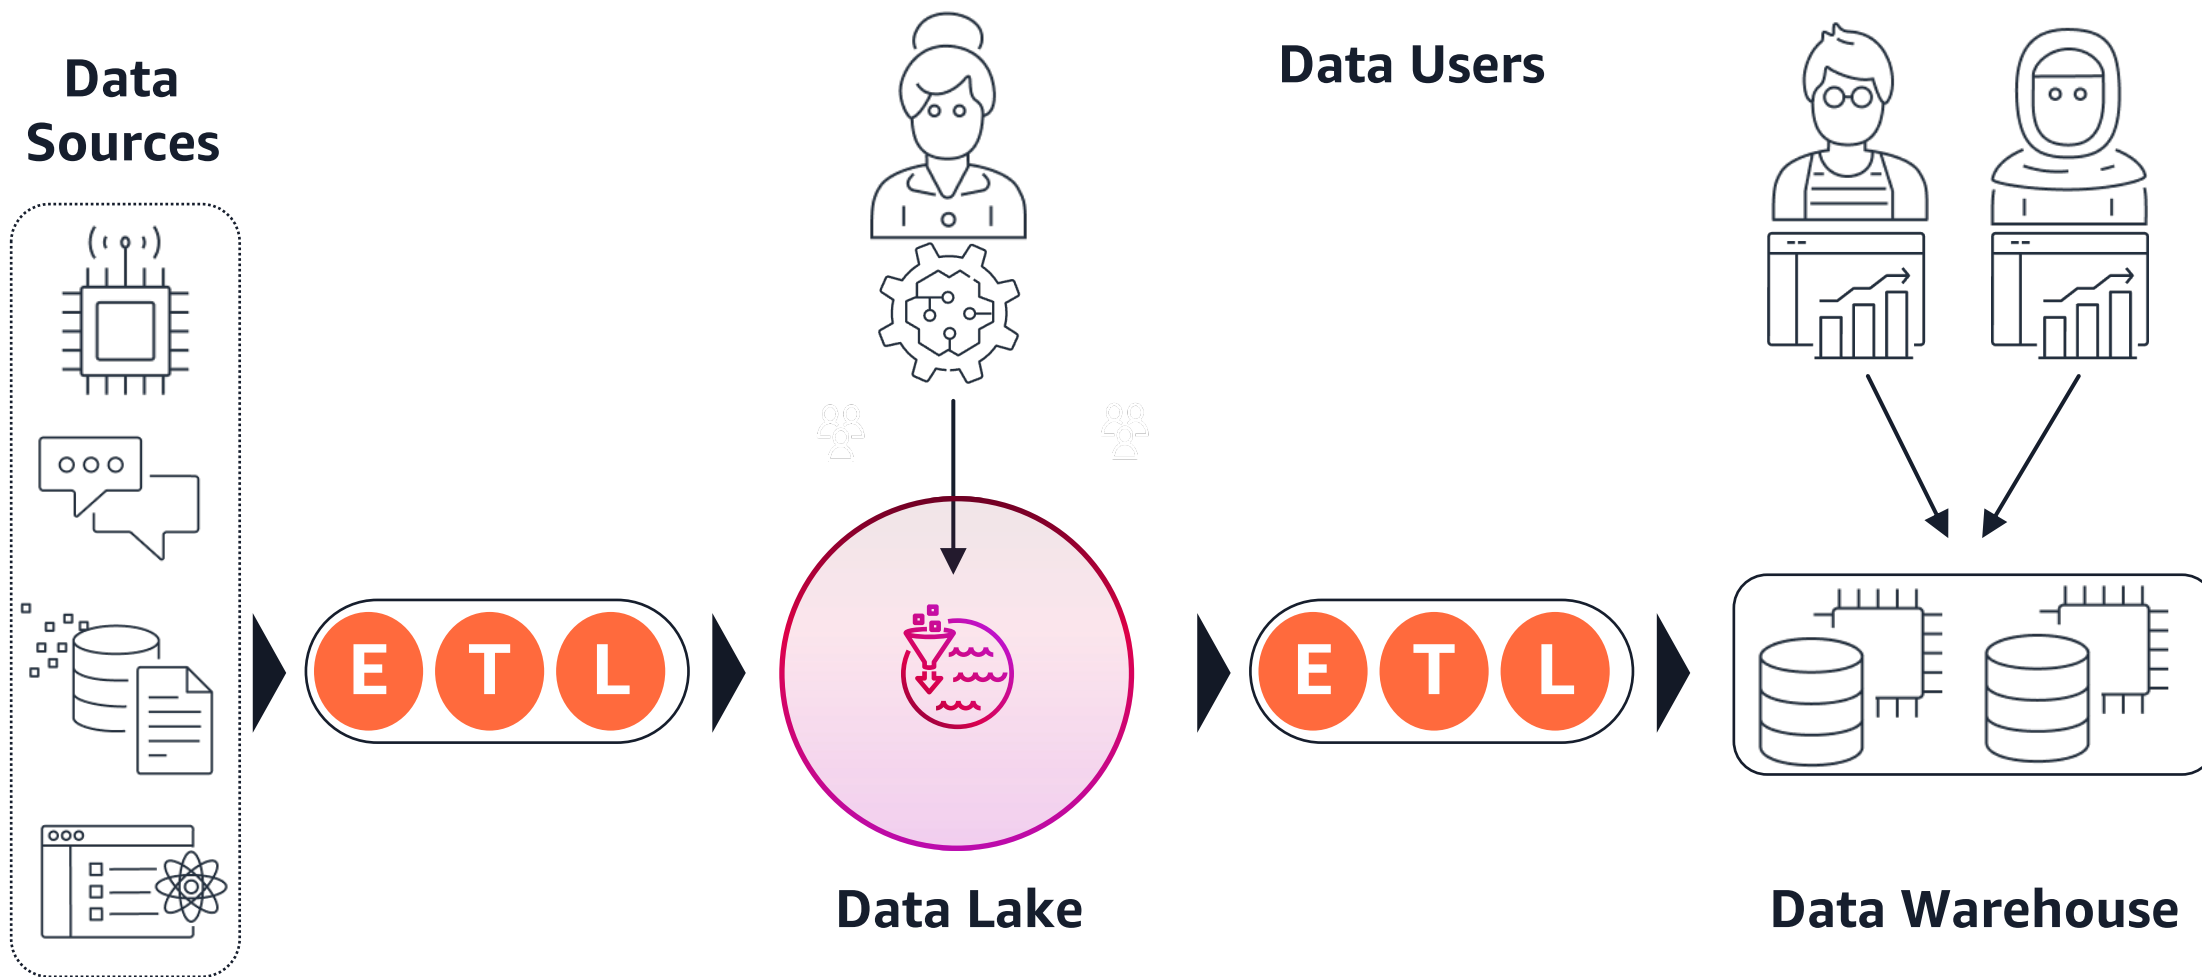


The rise of the Data Lake 2006+


Data Sources



The rise of the Data Lake 2006+



Challenges operating two independent systems



**No single
source of truth**

**Additional cost
and complexity**


**Error prone security
and governance**

Transactional Data Lakes


1 + 1 = Transactional Data Lake



Data Lake

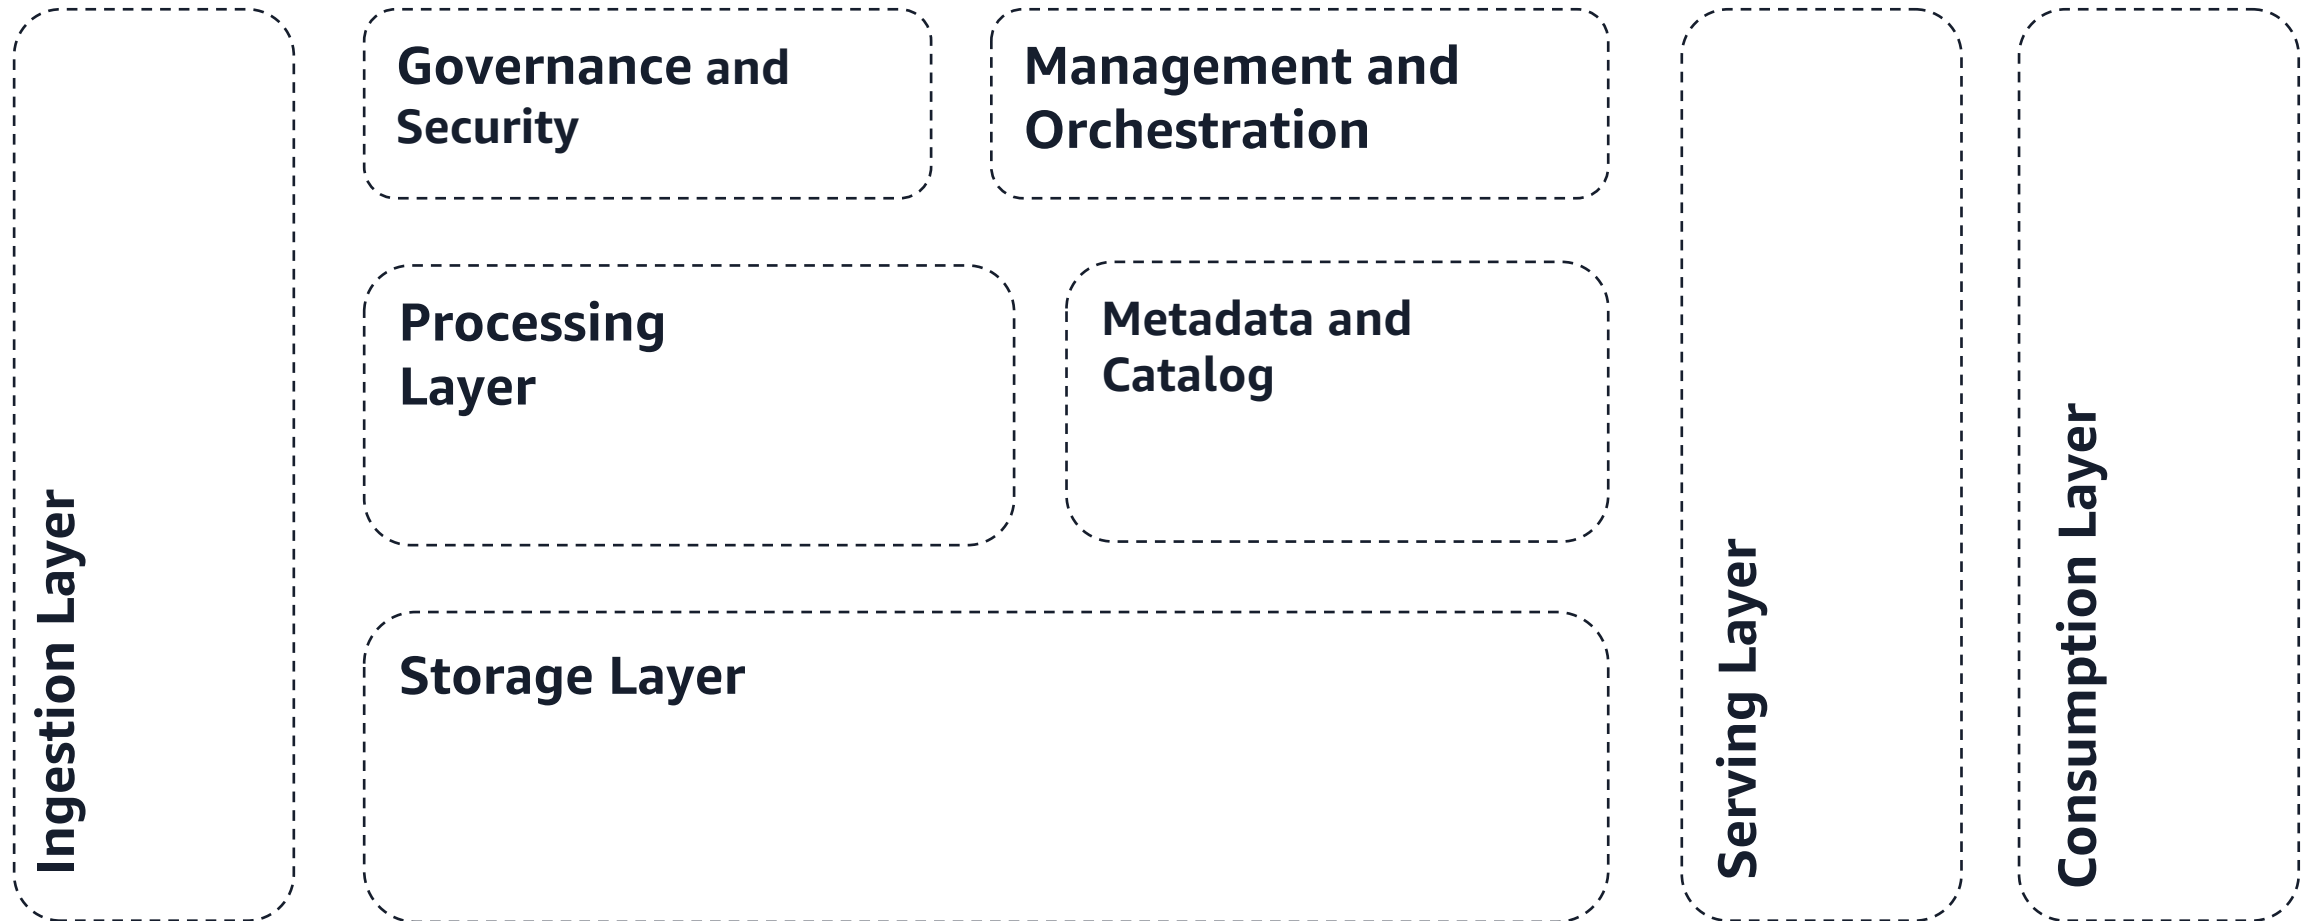


Data Warehouse

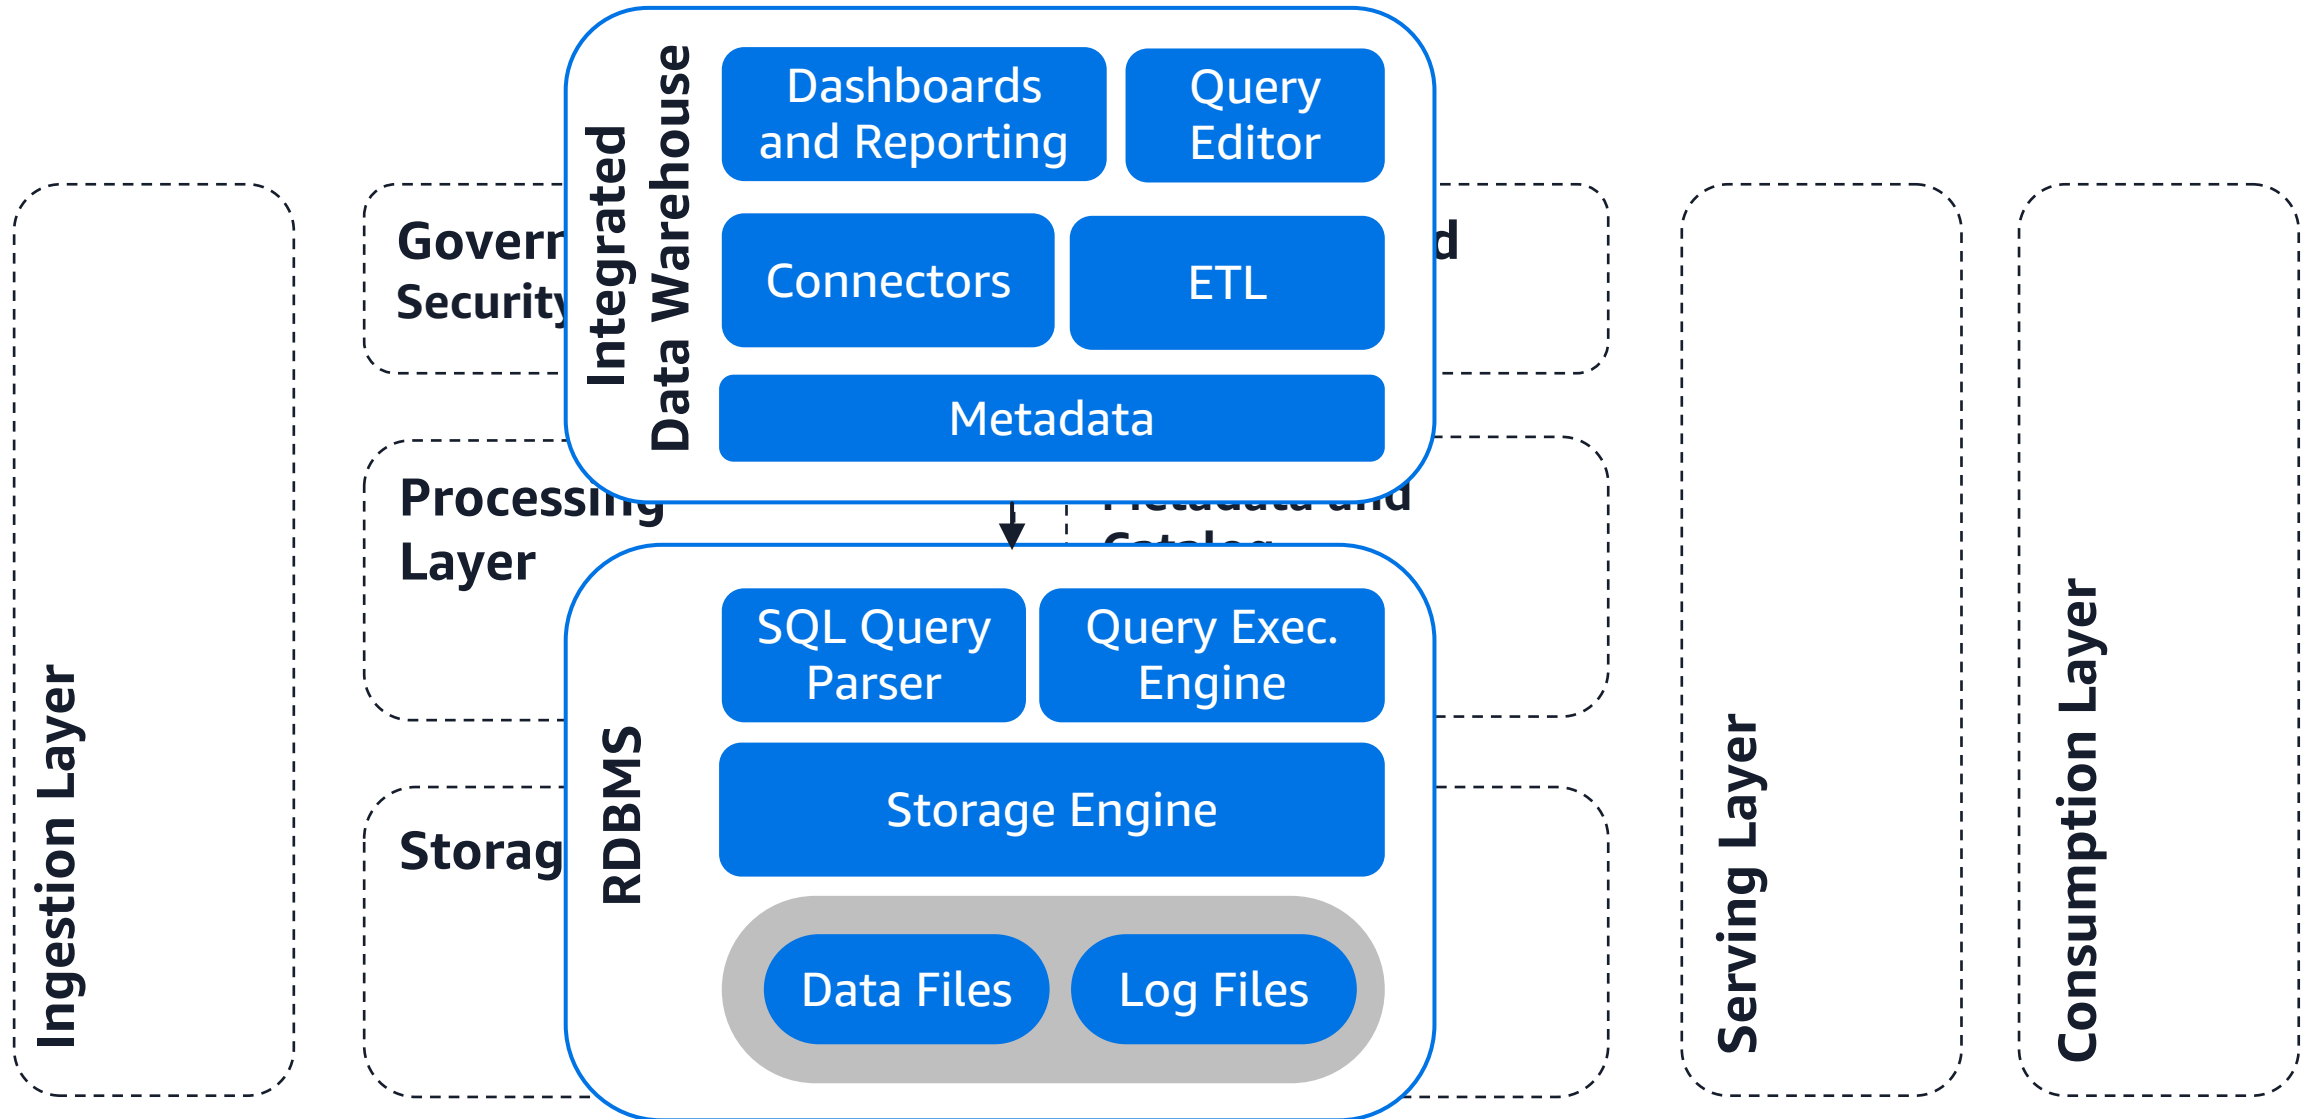


Transactional Data Lake

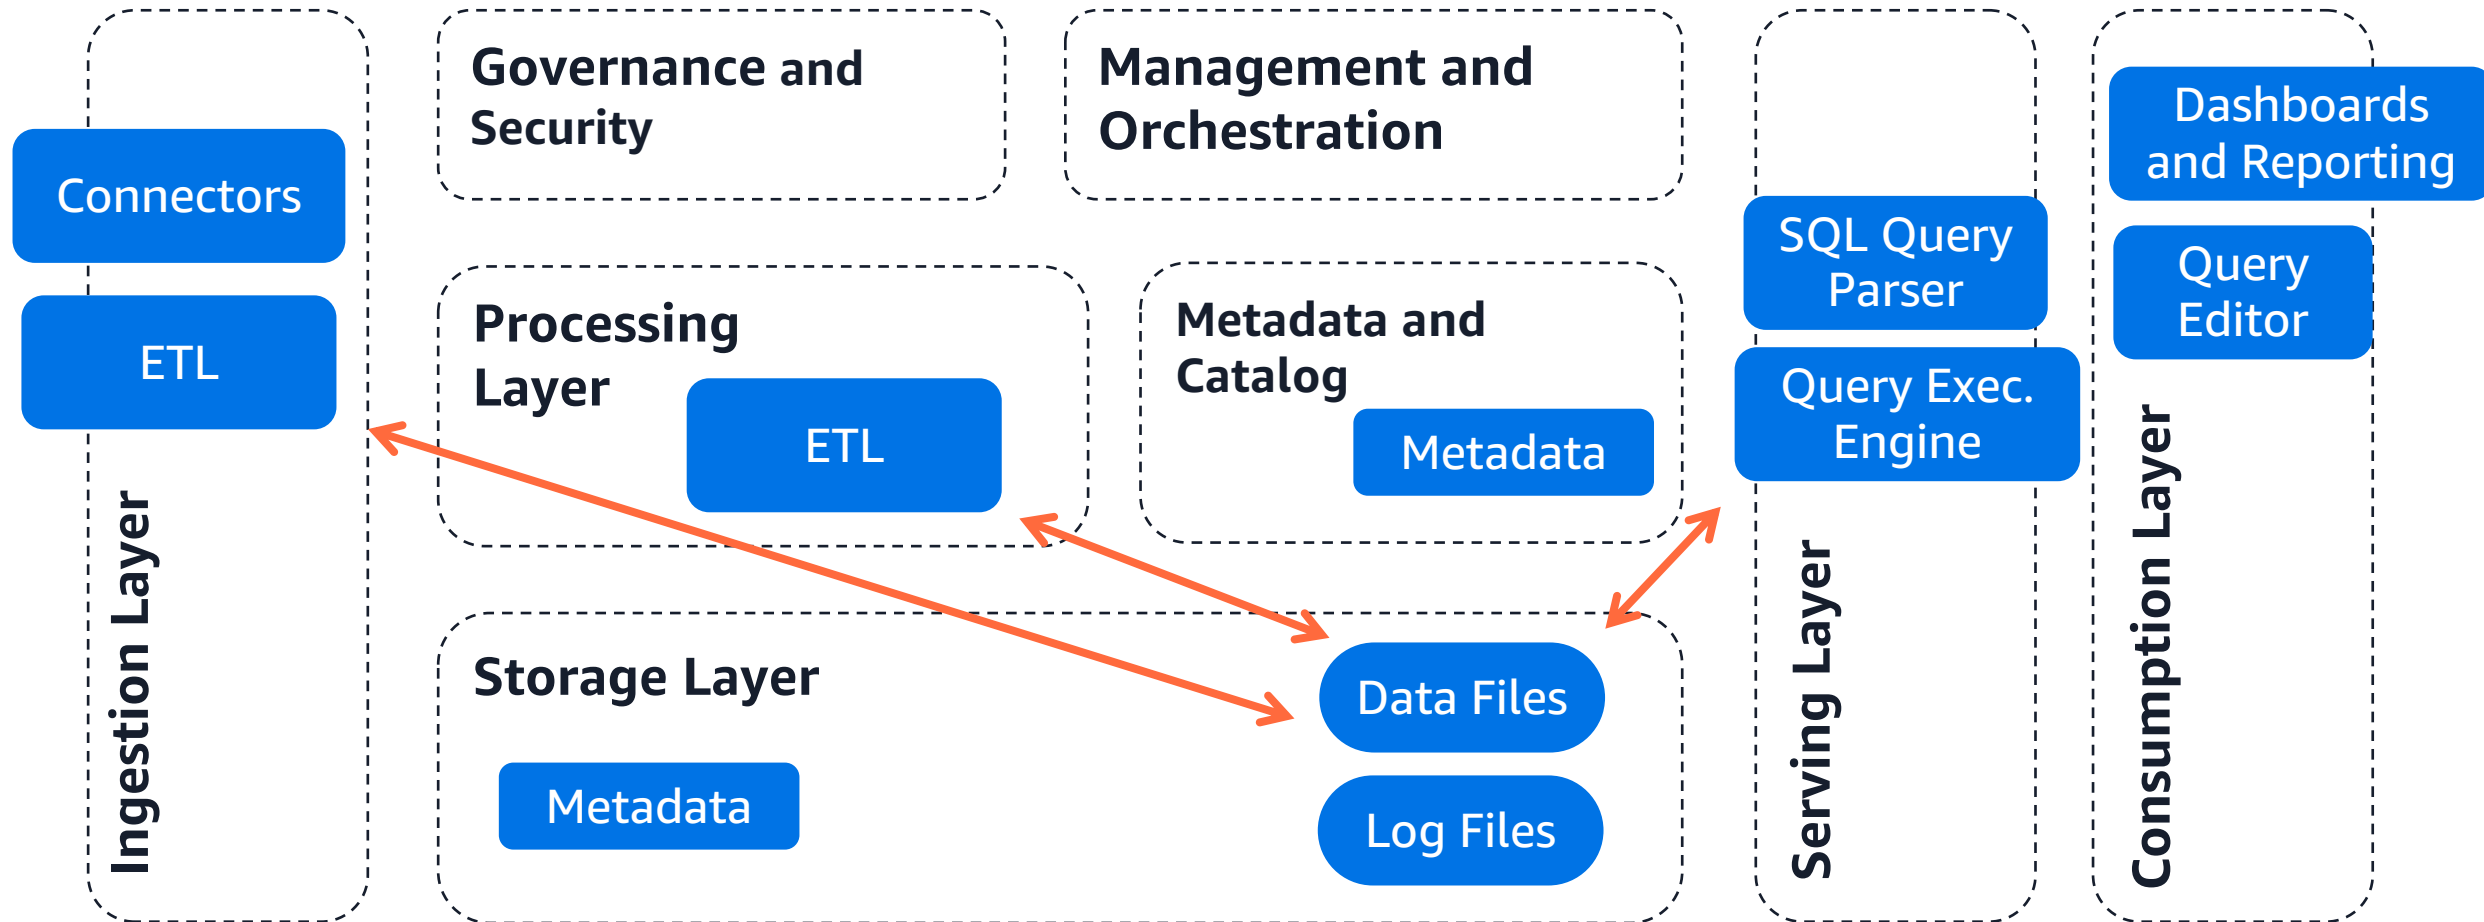
A typical data platform architecture



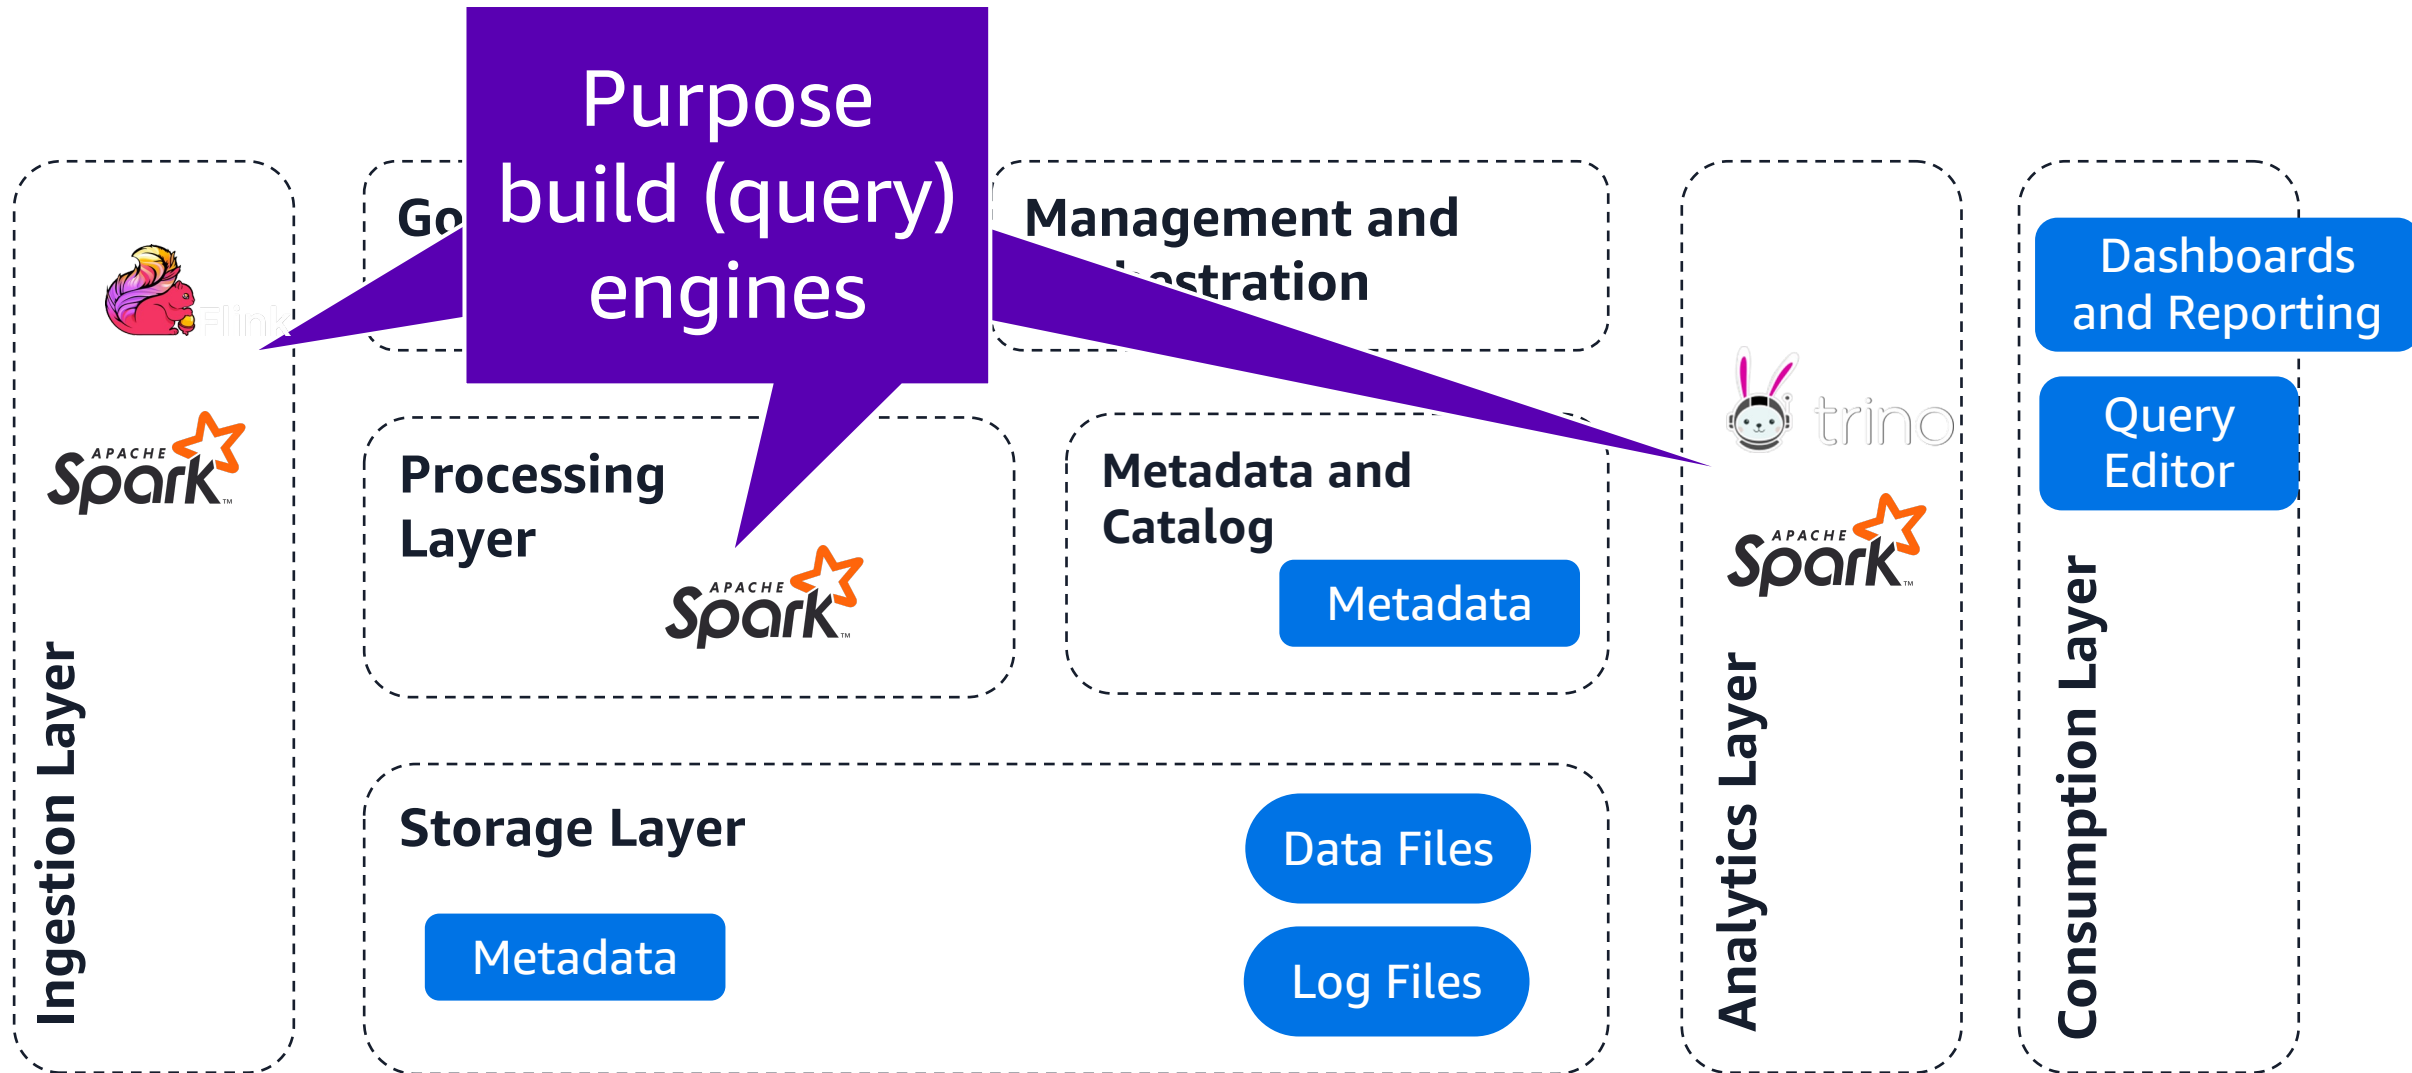
Turning the Data Warehouse inside out (1/2)



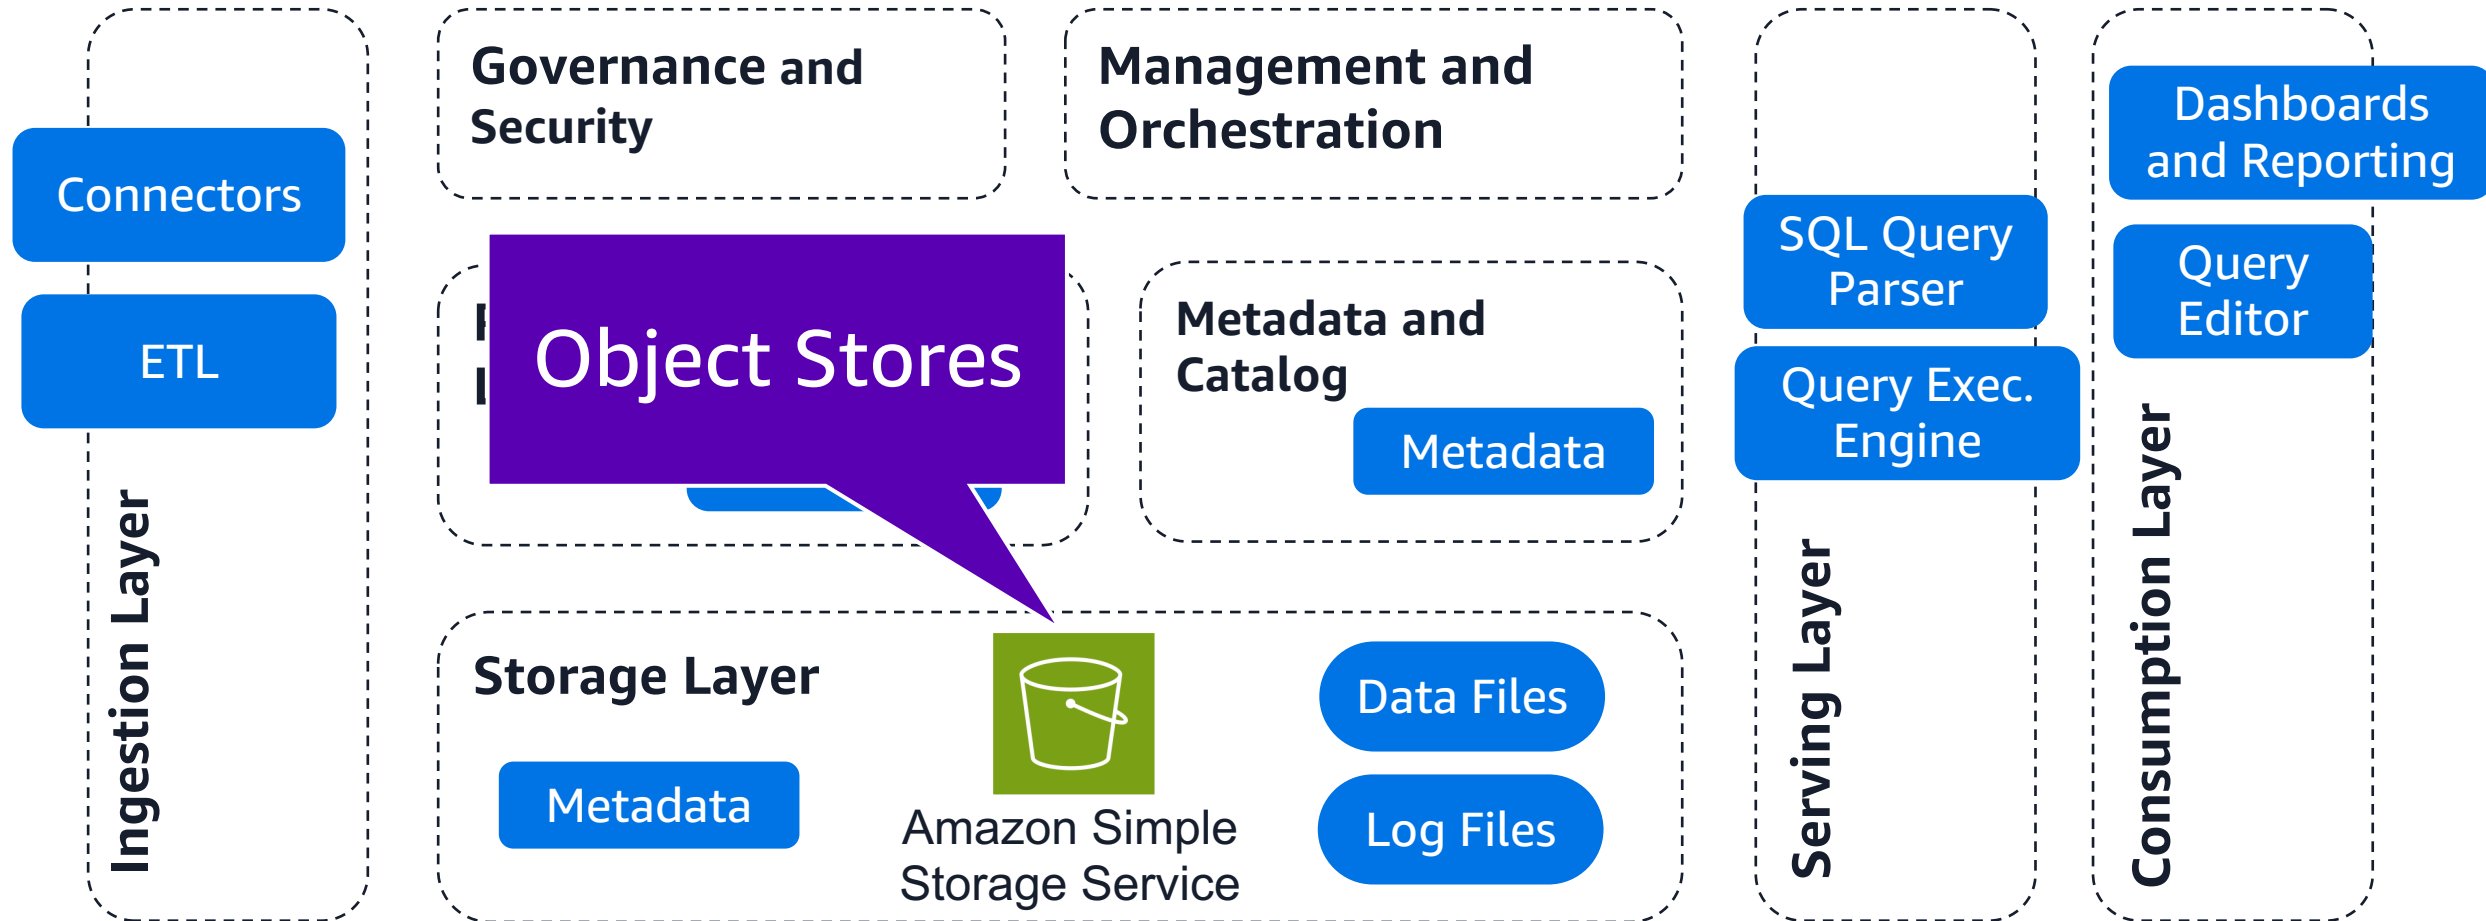
Turning the Data Warehouse inside out (2/2)



Powering the Transactional Data Lakes




Powering the Transactional Data Lake



Powered by object stores



Amazon Simple
Storage Service



infinite storage

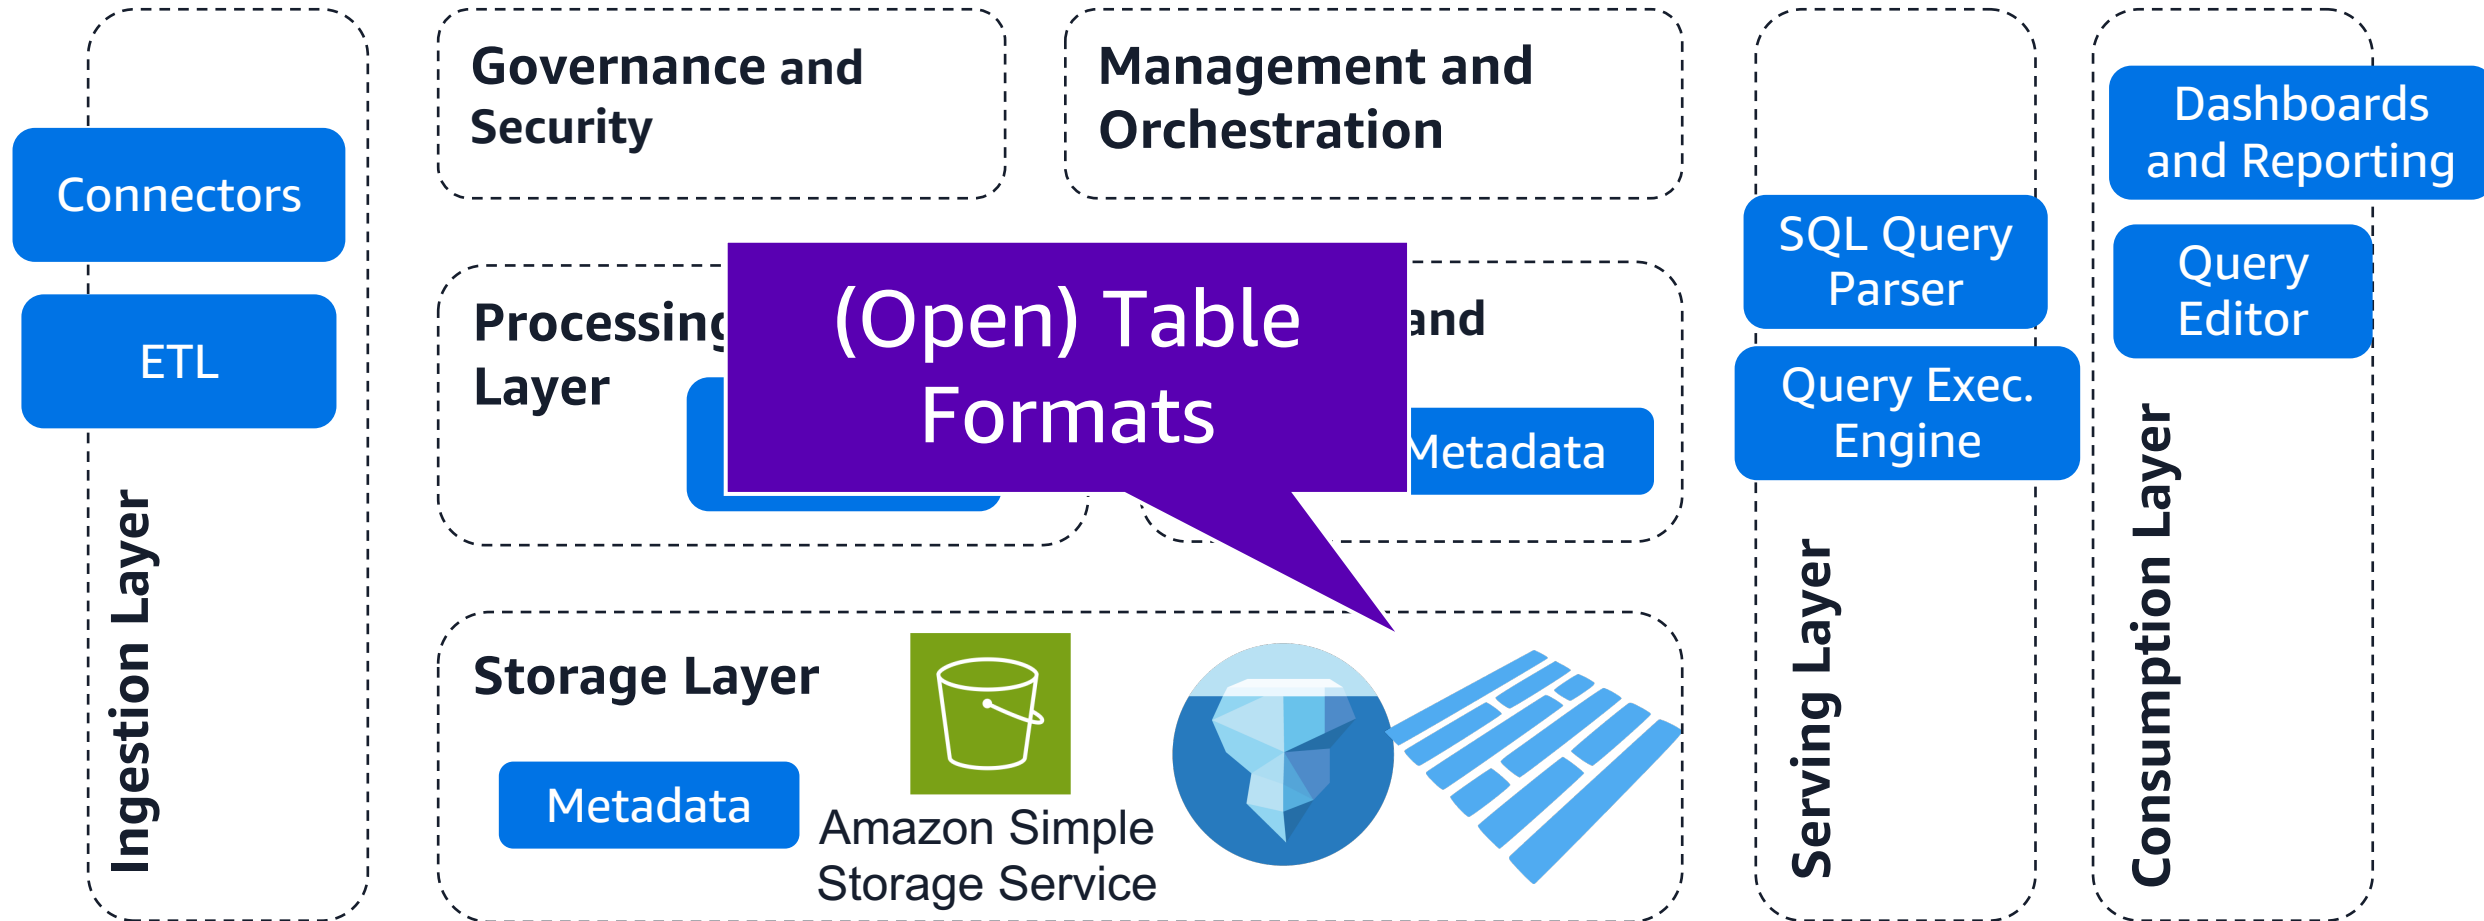


11 nines
of durability




99.99% availability

Powering the Transactional Data Lake



Backed by a broad choice of file formats



Common file formats

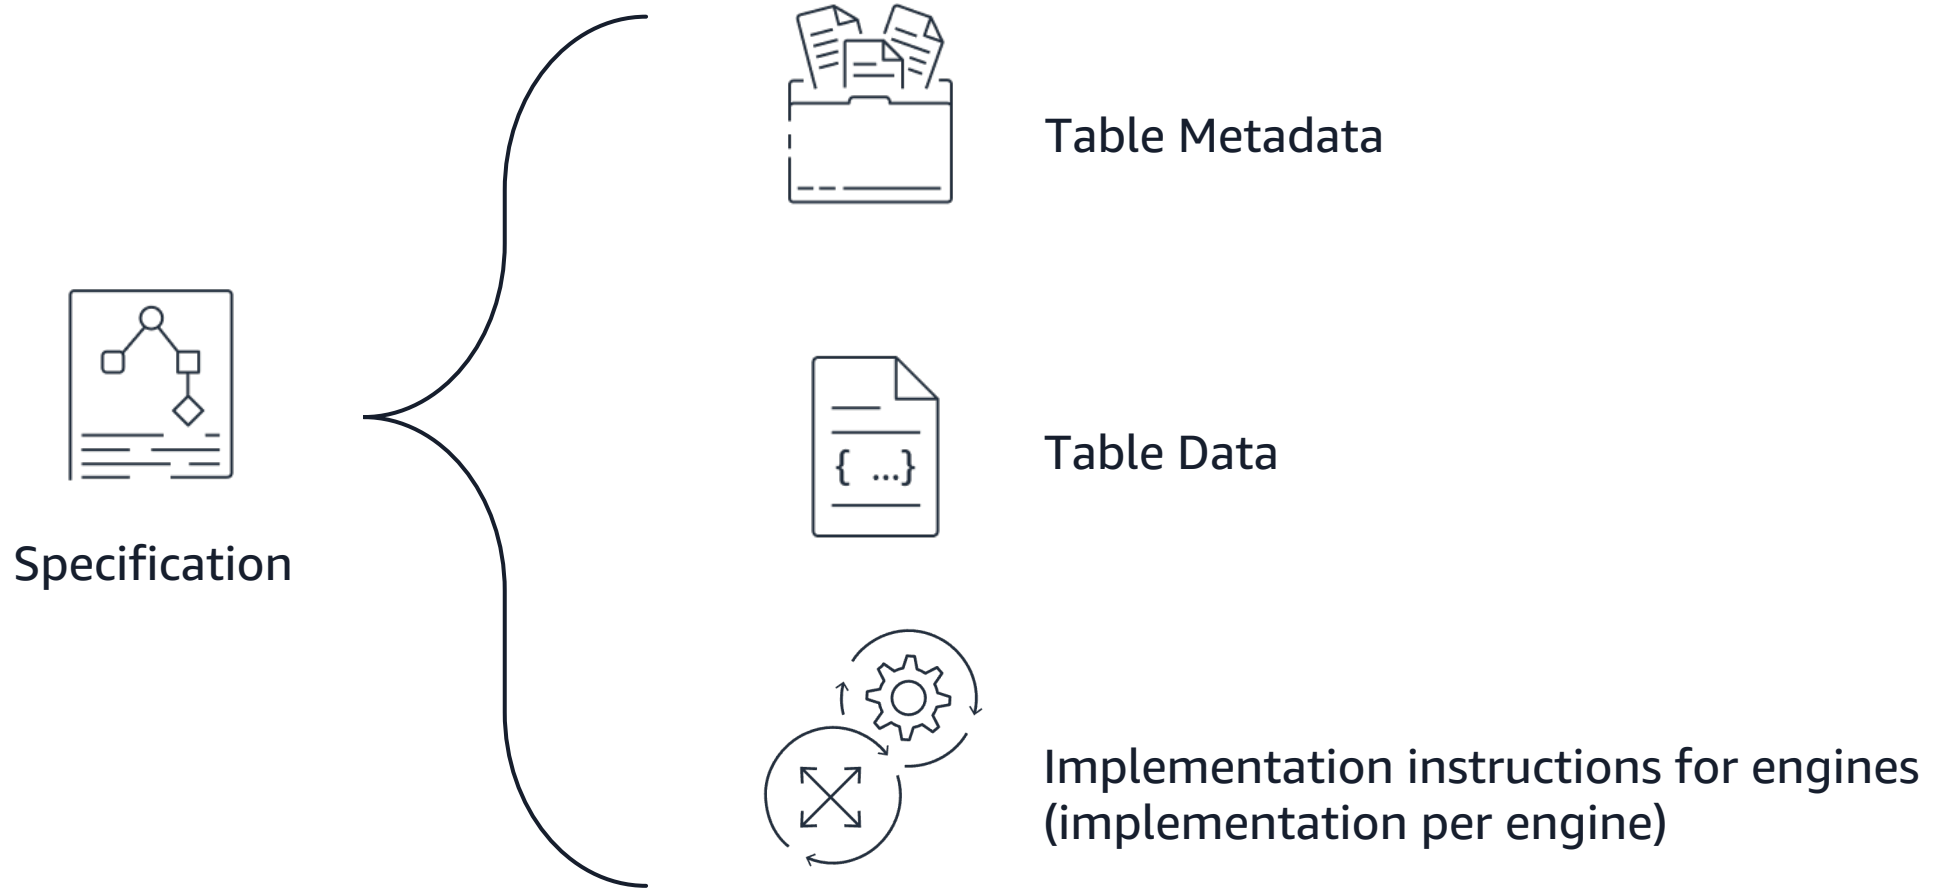


Optimized file formats



Table file formats

Open Table formats




Introduction to Apache Iceberg

Iceberg Ecosystem

Producers/Writers



ICEBERG Tables



Consumers/Readers



Iceberg Features

- **Delete, update, and upsert** support using copy-on-write or merge-on-read approaches
- **Fast scan planning & advanced filtering** using table metadata for with partition and column-level stats
- **Full schema evolution** supports add, drop, update, or rename, and has no side-effects
- **Partition layout evolution** can update the layout of a table as data volume or query patterns change
- **Time travel** enables reproducible queries that use exactly the same table snapshot, or lets users easily examine changes
- **Version rollback** allows users to quickly correct problems by resetting tables to a good state
- **Snapshot isolation** where table changes are atomic and readers never see partial or uncommitted changes
- **Multiple concurrent writers** use optimistic concurrency and will retry to ensure that compatible updates succeed, even when writes conflict
- ...

Iceberg Table Format Overview

- Iceberg tables are composed by 3 layers:
- Iceberg Catalog
- Metadata Layer
- Data Layer

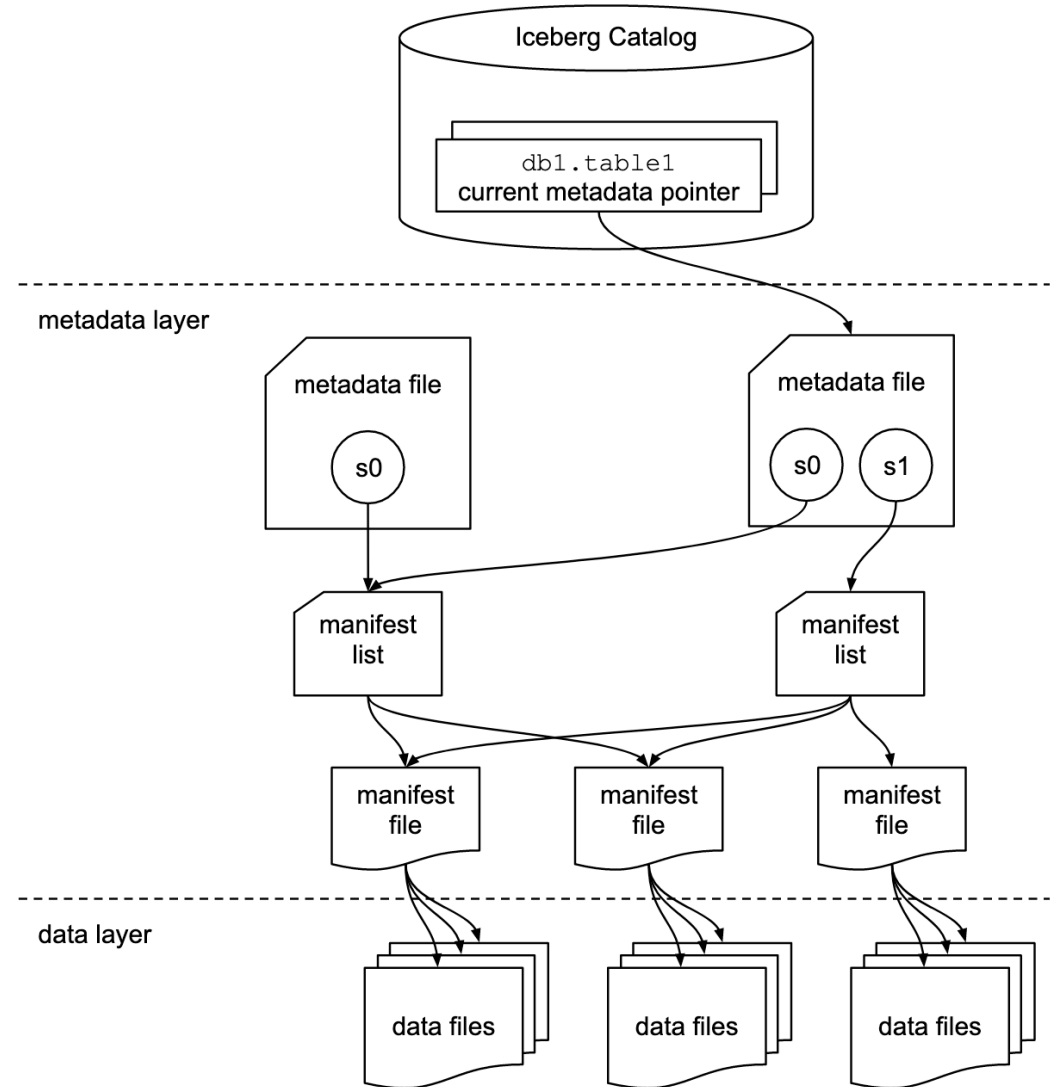


Table Layout - Catalog

Iceberg catalog stores the metadata pointer for the current table metadata file and supports atomic updates

Metadata pointer maps the table name to its current metadata file

When a **SELECT query** is reading an Iceberg table, the query engine first goes to the Iceberg catalog, then retrieves the entry of the location of the current metadata file

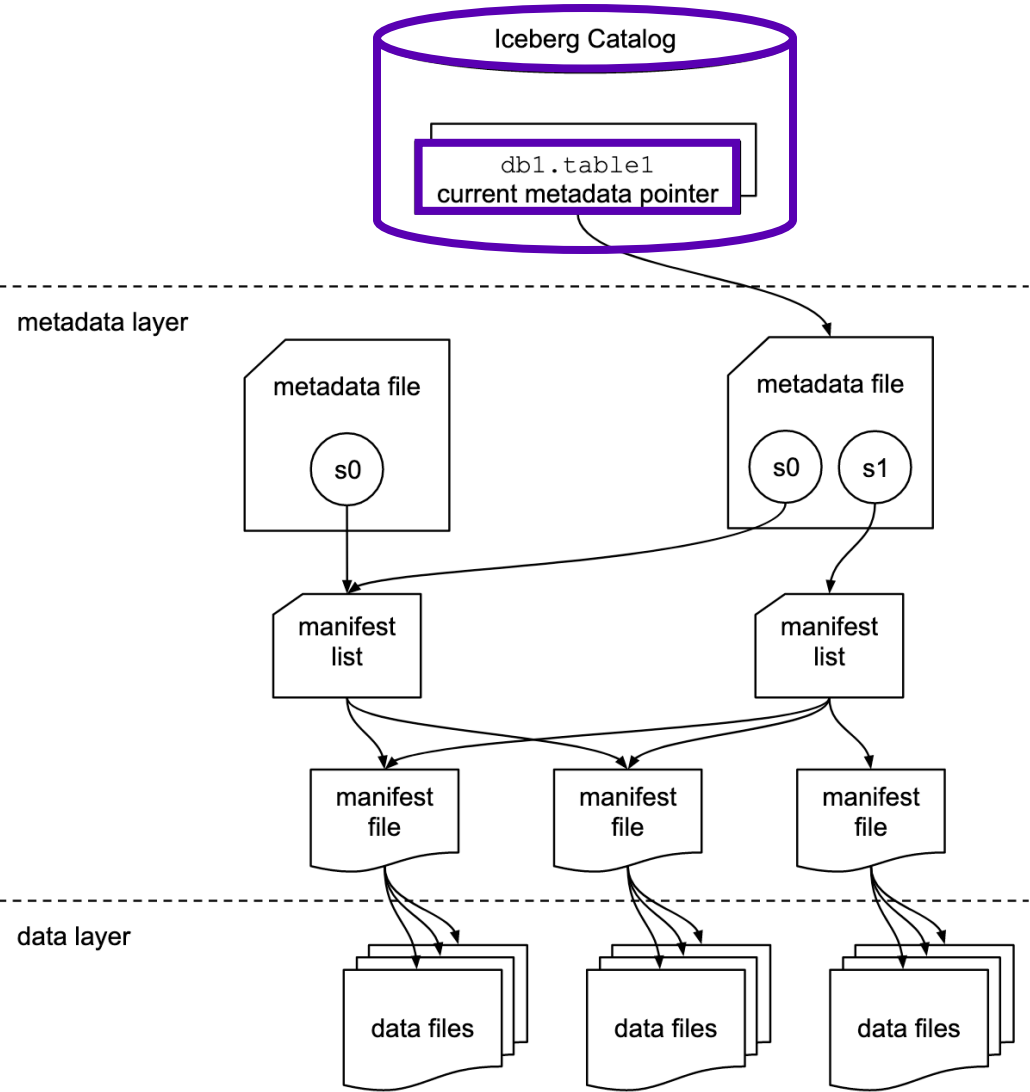


Table Layout - Metadata file

- **Metadata file** stores the table state. Any change creates a new metadata file replacing the old one with an **atomic swap**

- The table metadata file tracks **schema**, **partitioning**, **properties**, and **snapshots** of the table contents

- A **snapshot** represents the state of a table at some time and is used to access table's data files

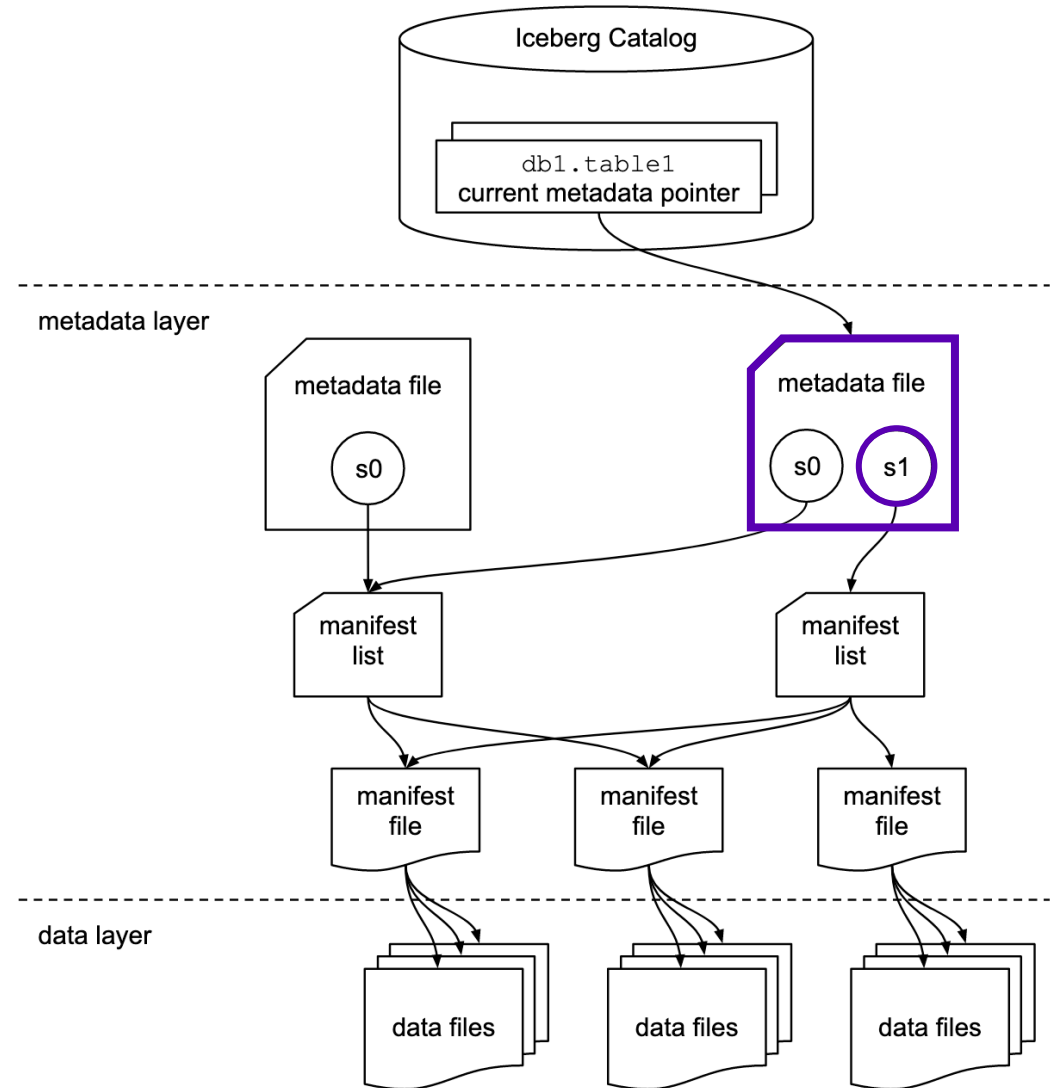


Table Layout - Metadata file

```
▼ snapshots:
  ▼ 0:
    sequence-number: 1
    snapshot-id: 8781292425728477000
    timestamp-ms: 1653321405356
    summary: {...}
    ▼ manifest-list: "s3:// [redacted] /warehouse/"
    schema-id: 0
  ▼ 1:
    sequence-number: 2
    snapshot-id: 4307640767686990000
    parent-snapshot-id: 8781292425728477000
    timestamp-ms: 1653321431440
    summary: {...}
    ▼ manifest-list: "s3:// [redacted] /warehouse/"
    schema-id: 0
  ▼ 2:
    sequence-number: 3
    snapshot-id: 3081751970069754400
    parent-snapshot-id: 4307640767686990000
    timestamp-ms: 1653321475470
    summary: {...}
    ▼ manifest-list: "s3:// [redacted] /warehouse/"
    schema-id: 0
```

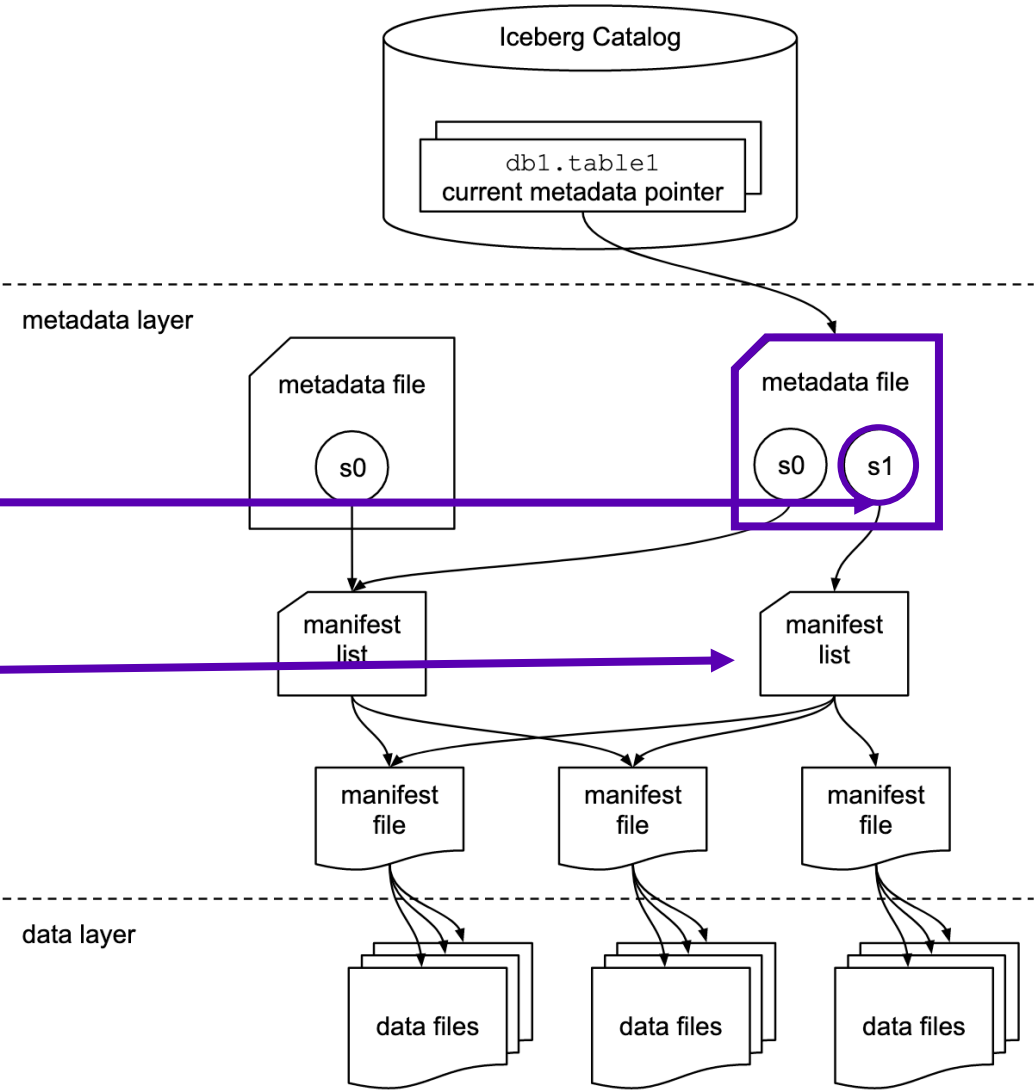


Table Layout - Manifest list file

```
{  
  "manifest_path": "s3://.../...6e8107509-m0.avro",  
  "added_snapshot_id": 8781292425728477000,  
  (...)  
},  
  
{  
  "manifest_path": "s3://.../...6678f597-m0.avro",  
  "added_snapshot_id": 3081751970069754400,  
  (...)  
},  
  
{  
  "manifest_path": "s3://.../....ejebe38-m0 avro",  
  "added_snapshot_id": S1,  
  (...)  
},  
  
{  
  "manifest_path": "s3://.../...c10be38-m1.avro",  
  "added_snapshot_id": S1,  
  (...)  
},  
...
```

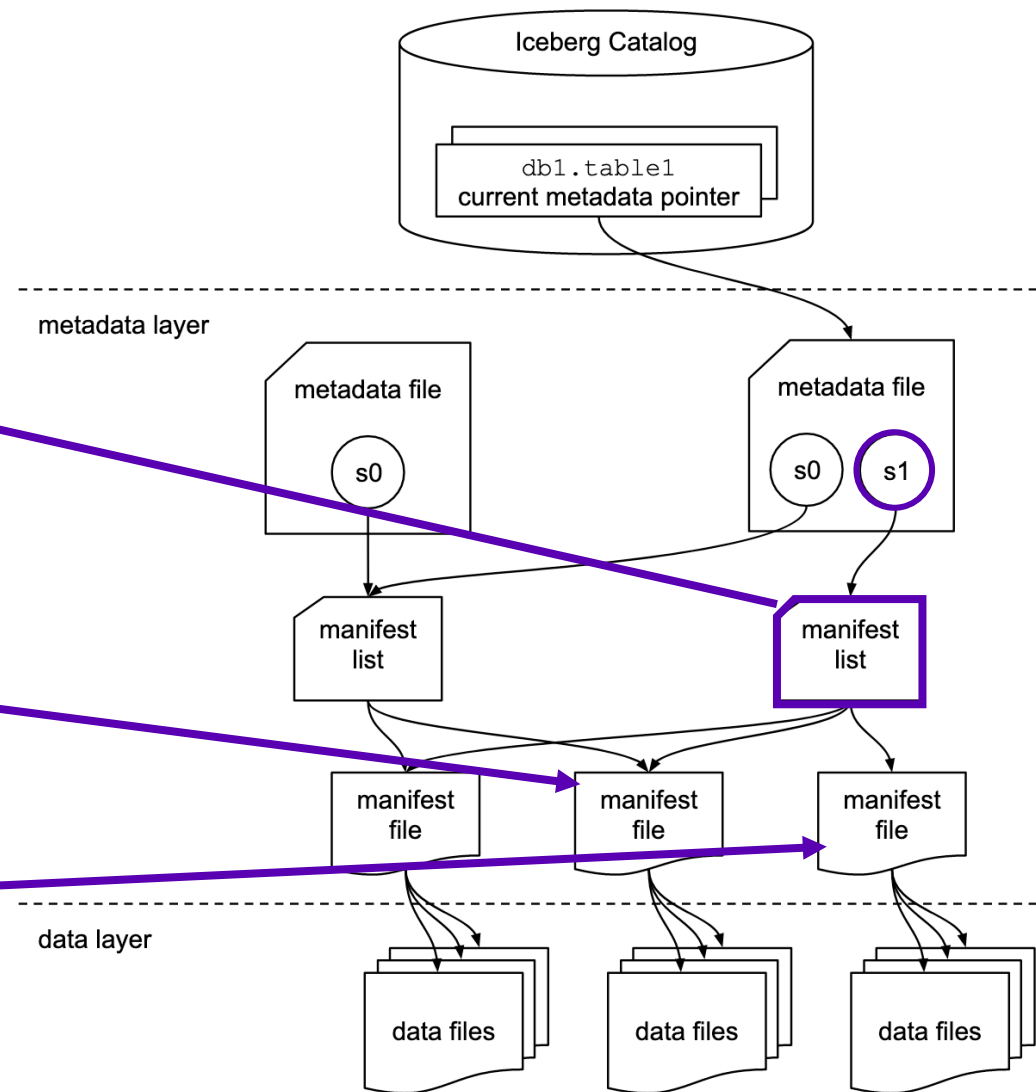


Table Maintenance

Recommended Maintenance

- Expire Snapshots
- Remove old metadata files
- Delete orphan files




Optional Maintenance

- Compaction of data files
- Compaction of manifest files


Transactional Data Lake on AWS

Transactional Data Lake reference architecture

Data Ingestion



Data Serving



Recap – What have we learned

Transactional Data Lakes combine the best of Data Lakes and Data Warehouses into one system, whilst avoiding their downsides.


Open Table Formats bring back ACID (Atomicity, Consistency, Isolation, Durability) transactions and enable advanced query optimization through extended metadata.

Apache Iceberg is an Open Table Format, widely supported inside and outside of AWS, keeping you in control of your data and giving you choice selecting the right engine for the job.

Thank you!

Norman Weisenburger

norwei@amazon.de

 [in /in/normalerweise](https://www.linkedin.com/company/norwei/)

Feedback



<https://pulse.aws/survey/ARYOGOIM>